Data Wrangling Part II
○○○○○○○○○○○

Data Cleaning
○○○

Descriptive Analytics
○○○○○

Class 5 Data Wrangling with R (Part II)

Dr Wei Miao

UCL School of Management

October 18, 2023

Section 1

Data Wrangling Part II

Data Wrangling Part II

- Select rows (`filter`)

- Sort rows (`arrange`)

- Select columns (`select`)

- Generate new columns (`mutate`)

- **Group aggregation (`group_by`): compute statistics for each group**

- **Merge datasets (`join`): combine datasets from different sources**

Aggregation by Groups: `group_by`

- `group_by()` allows us to aggregate data by group and compute statistics for each group

```
1   # group by marital status
2   data_demo %>%
3     group_by(Marital_Status)
```

- Internally, the dataset is already grouped based on the specified variable(s).
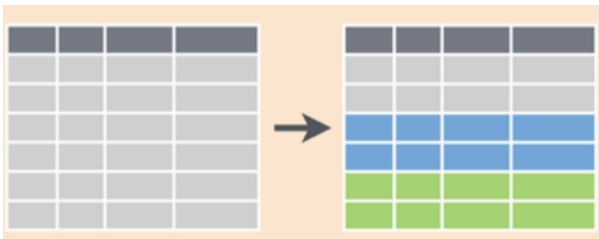
Aggregation by Groups: `group_by() %>% summarise()`

- After aggregating data, we can use `summarise()` to compute group-specific statistics for us.
  - Similar to `mutate()` in generating new variables
  - Different from `mutate()` in that the new variable is computed based on groups.

```
1  # compute the average income for each marital status group
2  data_demo %>%
3    group_by(Marital_Status) %>%
4    summarise(avg_income = mean(Income,na.rm = T)) %>%
5    ungroup()
```

Marital_Status	avg_income
Alone	43789.00
Divorced	53300.63
Married	51776.36
Single	51091.16
Together	52609.59
Widow	56158.90

- What if you replace `summarise()` with `mutate()`?

Aggregation by Groups: `group_by()` Multiple Groups

- We can have multiple group variables for `group_by`, such as computing average income for each marital status, education combination

```
# compute the average income for each marital, education group
data_demo %>%
  group_by(Marital_Status,Education) %>%
  summarise(avg_income = mean(Income,na.rm = T)) %>%
  ungroup() %>%
  head(5)
```

Marital_Status	Education	avg_income
Alone	Graduation	34176
Alone	Master	61331
Alone	PhD	35860
Divorced	2n Cycle	49345
Divorced	Basic	9548

Consolidate Multiple Data Frames

- When consolidating multiple data frames, we have 4 types of joining methods.
- `left_join()` handles most data join situations, which we will focus on today.

`left_join()`

- `left_join` keeps everything from the **left data frame** and matches as much as it can from the right data frame based on the chosen IDs.
  - All IDs **in the left data frame** will be retained
  - If a match can be found, value from the right data frame will be filled in
  - If a match cannot be found, a missing value will be returned

```
1  df_left %>%
2    left_join(df_right, by = c('ID' = 'ID') )
```
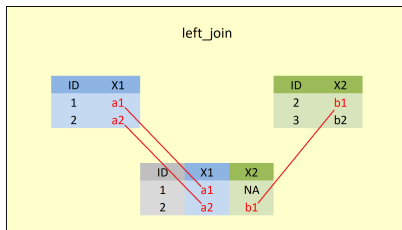
Caveats for doing `left_join()`

- We can do 1:1, or M:1 left_joins

- **Never** do 1:M or M:M left_joins

demographics: 1				Transaction data: M		
ID	age	income		ID	transaction	revenue
1	15	3000		1	1	12
2	18	1000		1	2	2
				2	1	10
				2	2	5
				2	3	10

inner_join() (optional)

- inner_join only keeps the observations that appear in both data frames
  - Only common IDs **in both data frames** will be retained
  - If a match can be found, values will be filled in from both data frames

```
1  # Method 1 without pipe operator
2  inner_join(df_left, df_right, by = 'ID')
3  # Method 2 with pipe operator
4  df_left %>%
5    inner_join(df_right, by = 'ID')
6  # Method 3: order of data frames should not matter. Why?
7  df_right %>%
8    inner_join(df_left, by = 'ID')
```
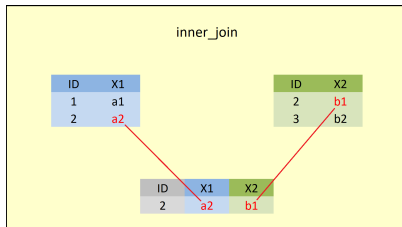
full_join() (optional)

- full_join keeps all observations from both data frames
  - All IDs **in either data frames** will be retained
  - If a match can be found, values will be filled in from both data frames

```
1  # Method 1 without pipe operator
2  full_join(df_left, df_right, by = 'ID')
3  # Method 2 with pipe operator
4  df_left %>%
5     full_join(df_right, by = 'ID')
6  # Method 3: order of data frames should not matter. Why?
7  df_right %>%
8     full_join(df_left, by = 'ID')
```

Data Wrangling Part II
0000000000

Data Cleaning
●○○

Descriptive Analytics
○○○○○
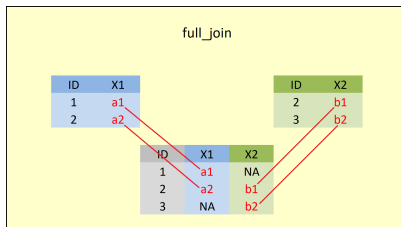
Section 2

Data Cleaning

Missing Values

- In R, missing values are represented by the symbol NA (i.e., *not available*).

- Most statistical models cannot handle missing values, so we need to deal with them in R.

  - Few missing values: remove them from analysis.

  - Many missing values: need to replace them with appropriate values: mean/median/imputation

Data Wrangling Part II
○○○○○○○○○○○

Data Cleaning
○○●

Descriptive Analytics
○○○○○

Outliers

- Sometimes, due to data collection errors, we may have abnormal observations in the data, such as unusually large and small values

- Winsorization is a common way to deal with outliers
  - Remove top 1% and bottom 1% observations

Section 3

Descriptive Analytics

Two Major Tasks of Descriptive Analytics

- You can think of descriptive analytics as **creating a dashboard** to display the key information you would like to know for your business.

1. **Describe data depending on your business purposes**
   - "How much do our customers spend each month on average?"
   - "What percentage of our customers are unprofitable?"
   - "What is the difference between the retention rates of men and women?"

2. **Make statistical inferences from data**
   - "Based on our sample, does the difference between the spending of men and women indicate that men and women respond differently in the customer base at large?"
   - "Based on our sample, can we conclude that customers who sign up for online banking are more profitable than customers who do not?"
   - "Based on our test mailing, can we conclude that ad-copy A works better than ad-copy B?"

Summary Statistics

- **Summary statistics** are used to summarize a set of observations, in order to communicate the largest amount of information as simply as possible.

- There are two main types of summary statistics used in evaluation:
  - **measures of central tendency**: mean, the median, 25 percentile, 75 percentile, the mode, etc.
  - **measures of dispersion:** range and standard deviation.

- It's important to include summary statistics table in your dissertation before any statistical analysis!

Data Wrangling Part II
0000000000

Data Cleaning
000

Descriptive Analytics
0000●0

Summary Statistics with R

- In R, a nice package to report summary statistics is `modelsummary`.

- `datasummary_skim()` is a shortcut to conduct basic summary statistics

- For more features, refer to the package tutorial here

```
1  pacman::p_load(modelsummary)
2  ## Summary statistics for numeric variables
3  data_demo %>%
4    datasummary_skim(type = "numeric")
5
6  ## Summary statistics for categorical variables
7  data_demo %>%
8    datasummary_skim(type = "categorical")
```

Data Wrangling Part II
0000000000

Data Cleaning
000

Descriptive Analytics
00000●

Case Study: Preliminary Customer Analysis

- Let's solve the preliminary customer analysis case together in class!