

Class 5 Data Wrangling with R (Part II)

Dr Wei Miao

UCL School of Management

Thu, Oct 20 2022

Section 1

Data Wrangling

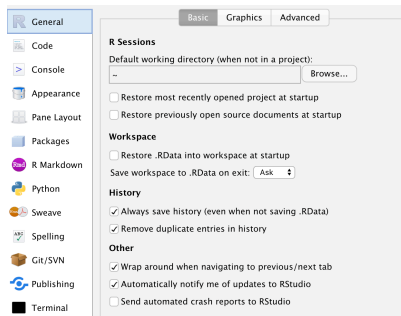
R Tips: More Convenient Package Management Using pacman

```
1  pacman::p_load(dplyr,ggplot2)
```

- Please install `pacman` on your RStudio
- `pacman`'s functionality
 - Load all packages stated in the parantheses, seperated by commas
 - If the package is not downloaded yet, download it, and then load it
- R tip: if you want to use a function without loading the whole package, you can use two colons to call the function: `package::function`

R Tips: Managing Objects in the RStudio Environment

- Best practice is to not save any objects once you close your RStudio session



- `rm(list = ls())` is the command to remove everything in the current environment
 - `ls()` is a function that returns the list of all objects in the current environment
 - `rm(list =)` removes any objects passed to `list` argument

Recap: filter(), arrange(), and mutate()

- `filter(dataset, criteria)`: pick observations by their values

Subset Observations (Rows)



- `arrange(dataset, variable)`: reorder the rows
- `mutate(dataset, newvariable =)`: create new variables with functions of existing variables

Make New Variables



- Imagine a factory with different machines placed along a belt. Each machine is a **dplyr function** that performs a data cleaning step, like filtering or arranging data.
- The pipe therefore works like a conveyor belt, passing the output of one machine to another for further processing.

input %>% functionName(, ...)

↕

functionName(input, ...)

Pipe Operator

- The pipe has a huge advantage over any other method of processing data in R or Python: It makes data wrangling processes easy to read. If we read `%>%` as “then”, the code will be very easy to interpret as a set of instructions in plain English:

```

1 mtcars %>%
2   filter(cyl>=5) %>%
3   mutate(sqrt_wt=sqrt(wt)) %>%
4   filter(sqrt_wt>1.5) %>%
5   arrange(hp)
6 ## can go on and on

```

- 1 take the mtcars data, THEN
- 2 find cars cyl ≥ 5, THEN
- 3 mutate a new variable sqrt_wt, THEN
- 4 find cars sqrt_wt > 1.5 THEN
- 5 reorder all cars based on hp
- 6 chain more operations ...

Without Pipe Operator

- As a comparison, without using pipe operators, the previous data cleaning steps need to be done as follows. Overwriting our output dataframe `new_data` in every line is problematic.
 - First, doing this for a procedure with lots of steps isn't efficient and creates unnecessary repetition in the code.
 - Second, this repetition also makes it harder to identify exactly what is changing on each line in some cases.

```
1 new_data <- filter(mtcars, cyl >=5)
2 new_data <- mutate(new_data, sqrt_wt=sqrt(wt))
3 new_data <- filter(new_data,sqrt_wt>1.5)
4 new_data <- arrange(new_data,hp)
```


Select Variables: select

- `select()` can select variables into a smaller dataset.

```
1 # Select two columns: hp and cyl
2 mtcars%>%
3   select(hp, cyl) %>%
4   head()
```

	hp	cyl
Mazda RX4	110	6
Mazda RX4 Wag	110	6
Datsun 710	93	4
Hornet 4 Drive	110	6
Hornet Sportabout	175	8
Valiant	105	6

Aggregation by Groups: `group_by`

- `group_by()` allows us to aggregate data by group and compute statistics for each group

```
1 # group by cyl
2 mtcars %>%
3   group_by(cyl)
```

- Although nothing seemingly happens to the dataset, internally, the dataset is already grouped based on the specified variable(s).



Aggregation by Groups: `group_by()` + `summarise()`

- `summarise()` creates a new data frame after aggregating data. The final dataset
 - has one row for each pair of grouping variables (for each `cyl` value)
 - contains one column for each grouping variable (`cyl`)
 - contains one column for each new summarised variable (`avg_mpg`)

```
1 # compute the average mpg for each cyl group
2 mtcars %>%
3   group_by(cyl) %>% # group by cyl
4   summarise(avg_mpg = mean(mpg)) %>% # compute the average mpg
5   ungroup()
```

cyl	avg_mpg
4	26.66364
6	19.74286
8	15.10000

Aggregation by Groups: `group_by()` + `summarise()`

- We can have multiple group variables for `group_by`

```
1 # compute the average mpg for each cyl,vs group
2 mtcars %>%
3   group_by(cyl,vs) %>% # group by cyl
4   summarise(avg_mpg = mean(mpg)) %>% # compute the average mpg
5   ungroup()
```

Aggregation by Groups: `group_by()` + `mutate()`

- Try the following code by replacing `summarise()` with `mutate()`, what do you get now?

```
1 # compute the average mpg for each cyl,vs group
2 mtcars %>%
3   group_by(cyl,vs) %>% # group by cyl
4   mutate(avg_mpg = mean(mpg)) %>% # compute the average mpg
5   ungroup()
```

- A new column is added to the original dataset, the value of which is from the group-by aggregation.

Consolidate Multiple Data Frames

- When consolidating multiple data frames, we usually have 4 types of joining methods

Combine Data Sets

a		b	
x1	x2	x1	x3
A	1	A	T
B	2	B	F
C	3	D	T

+

=

Mutating Joins

x1	x2	x3
A	1	T
B	2	F
C	3	NA

dplyr::left_join(a, b, by = "x1")

Join matching rows from b to a.

x1	x3	x2
A	T	1
B	F	2
D	T	NA

dplyr::right_join(a, b, by = "x1")

Join matching rows from a to b.

x1	x2	x3
A	1	T
B	2	F

dplyr::inner_join(a, b, by = "x1")

Join data. Retain only rows in both sets.

x1	x2	x3
A	1	T
B	2	F
C	3	NA
D	NA	T


dplyr::full_join(a, b, by = "x1")

Join data. Retain all values, all rows.

left_join

- left_join keeps everything from the left data frame and matches as much as it can from the right data frame.
 - All IDs in the left data frame will be retained
 - If a match can be found, value from the right data frame will be filled in
 - If a match cannot be found, a missing value will be filled in


```
1 # Method 1 without pipe operator
2 left_join(df_left, df_right, by = 'ID')
3 # Method 2 with pipe operator
4 df_left %>%
5   left_join(df_right, by = 'ID')
```



inner_join

- `inner_join` only keeps the observations that appear in both data frames
 - Only common IDs in **both data frames** will be retained
 - If a match can be found, values will be filled in from both data frames


```
1 # Method 1 without pipe operator
2 inner_join(df_left, df_right, by = 'ID')
3 # Method 2 with pipe operator
4 df_left %>%
5   inner_join(df_right, by = 'ID')
6 # Method 3: order of data frames should not matter. Why?
7 df_right %>%
8   inner_join(df_left, by = 'ID')
```



full_join

- `full_join` keeps all observations from both data frames
 - All IDs in **either data frames** will be retained
 - If a match can be found, values will be filled in from both data frames

```
1 # Method 1 without pipe operator
2 full_join(df_left, df_right, by = 'ID')
3 # Method 2 with pipe operator
4 df_left %>%
5   full_join(df_right, by = 'ID')
6 # Method 3: order of data frames should not matter. Why?
7 df_right %>%
8   full_join(df_left, by = 'ID')
```



Section 2

Data Cleaning

Variable Types

- Non-metric
 - Categorical (gender, region, brand, religion)
 - Ordinal (Business Week rankings, NCAA rankings)
- Metric
 - Continuous (age, height, sales, rainfall)
- Different types of variables are handled in different ways in statistics
 - Can talk about an average age, but not an average color
 - Some statistical techniques only work with one type of variable
- We need to make sure the variables are of the correct data types. Or we may need to convert them to the correct types.
 - e.g., from character to date time using `lubridate` package

Missing Values

- In R, missing values are represented by the symbol NA (i.e., *not available*).
- Most statistical models cannot handle missing values, so we need to deal with them in R.
 - Few missing values: remove them from analysis.
 - Many missing values: need to replace them with appropriate values: mean/median/[imputation](#)

Outliers

- In statistics, an outlier is a data point that differs significantly from other observations.
 - Few outliers: remove them from analysis
 - Many outliers: [winsorize](#) data
 - **If the distribution of a variable is not normal distribution, we often log transform variables to mitigate outlier issues**

Section 3

Descriptive Analytics

Two Major Tasks of Descriptive Analytics

1 Describe data depending on your business purposes

- “How much do our customers spend each month on average?”
- “What percentage of our customers are unprofitable?”
- “What is the difference between the retention rates of men and women?”

2 Make statistical inferences from data

- “Based on our sample, does the difference between the spendings of men and women indicate that men and women respond differently in the customer base at large?”
- “Based on our sample, can we conclude that customers who sign up for online banking are more profitable than customers who do not?”
- “Based on our test mailing, can we conclude that ad-copy A works better than ad-copy B?”

Descriptive Analytics

- You can think of descriptive analytics as **creating a dashboard** to display the key information you would like to know for your business.

The screenshot shows the YouTube Studio Channel dashboard for a user named 'MeowMeow'. The interface is divided into several sections:

- Left Sidebar:** Contains navigation links for Dashboard (highlighted), Content, Playlists, Analytics, Comments, Subtitles, Copyright, Monetisation, Customisation, and Audio library. At the bottom are Settings and Send feedback.
- Header:** Includes the YouTube Studio logo, a search bar for the channel, and notification icons.
- Channel dashboard:**
 - Latest video performance:** Features a video thumbnail titled 'Shape of you ukulele cover'. Below it, a table shows performance metrics for the first 29 days: Ranking by views (1 of 5), Views (41), Impressions click-through rate (20.5%), and Average view duration (0:43). Links for 'GO TO VIDEO ANALYTICS' and 'SEE COMMENTS (1)' are provided.
 - Channel analytics:** Displays 'Current subscribers' as 6, with a note '+5 in last 28 days'. It includes a 'Summary' section for 'Views' (88) and 'Watch time (hours)' (1.4). A 'Top videos' section lists '恋爱的犀牛 Rhinoceros in love guitar cover' (2 views) and 'Hear Me Cry Guitar Fingerstyle' (1 view). A link 'GO TO CHANNEL ANALYTICS' is at the bottom.
 - Latest comments:** Shows channel comments that haven't been responded to, featuring a comment from 'Hannah Grace Ukulele' about a performance. A 'VIEW MORE' link is present.
 - Recent subscribers:** A section for subscribers from the last 90 days.
- Right Panel:** Titled 'News', it features an illustration of a smartphone with a mail icon and a message: 'FYI: Handles are coming to YouTube. We're making it easier for people to find and connect with you. You'll get an email and a notification in Studio when you can choose yours.' A 'LEARN MORE' link is at the bottom.

Summary Statistics

- In descriptive analytics, **summary statistics** are used to summarize a set of observations, in order to communicate the largest amount of information as simply as possible.
- There are two main types of summary statistics used in evaluation: **measures of central tendency** and **measures of dispersion**.
 - Measures of central tendency provide different versions of the average, including the mean, the median, 25 percentile, 75 percentile, the mode, etc.
 - Measures of dispersion provide information about how much variation there is in the data, including the range and the standard deviation.
- It's good to include summary statistics table in your dissertation before any statistical analysis!
 - Commonly reported summary statistics include *mean*, *standard deviation*, *number of observations*, *min*, *25 percentile*, *median*, *75 percentile*, and *max*.
 - Then describe the distribution of the variable, dispersion of the variable, etc.

Summary Statistics with R

- In R, a nice package to report summary statistics is `modelsummary`.
`datasummary_skim()` is a shortcut to conduct basic summary statistics
- For more features, refer to the package tutorial [here](#), especially `datasummary()` function.
 - `datasummary_skim()` is a special case of more general `datasummary()`, which outputs a pre-determined set of summary statistics

```
1  pacman::p_load(modelsummary)
2  mtcars %>%
3    datasummary_skim()
```

Correlation Matrix

- Correlation matrix helps us understand the co-movement of any two variables in the data
- `datasummary_correlation()` reports the pairwise correlation coefficient
- In general, in a statistical model, variables of high correlation should not be included together, which leads to instability

```
1 mtcars %>%  
2   datasummary_correlation()
```

Section 4

Preliminary Customer Analysis

Preliminary Customer Analysis

- Spend 30 min-ish to work on the case study in a group.
- There are 7 questions in total. At the end of the discussion, each group selects a group leader to answer one question.
- To show your codes, join the Zoom link under “Module Overview” on Moodle.
- The group that did not answer correctly needs to do a performance next week!

After-Class Exercise

- What percent of customers are single? Try alternative ways to do the calculation.
- Is the average total spending by responders and non-responders statistically different? Answer this question using a t-test.
- Is income and total spending correlated?
- Are PhDs more likely to respond to marketing offers than Graduation? Use a statistical test to answer the question. Is the result what you expected?
- What would be the other useful descriptive analytics you would like to know for Tesco?