

Class 17 Regression Discontinuity Design

Dr Wei Miao

UCL School of Management

November 29, 2023

Section 1

Natural Experiment

From RCTs to Secondary Data

- RCTs are the gold standard of causal inference: In an RCT, the treatment is randomized and hence uncorrelated with any confounding factors, i.e., $cov(X, \epsilon) = 0$
- In practice, however, it can be challenging to implement a perfect RCT.
 - 1 Crossover and spillover effects;
 - 2 Costly in terms of time and money
- Therefore, we may want to exploit causal effects from existing secondary data. Besides the **instrumental variable** method, we can also investigate **natural experiments**.

Comparison: RCT & Natural Experiment

i Natural Experiment

A **natural experiment** is an event in which individuals are exposed to the experimental conditions that are determined by **nature** or **exogenous factors beyond researchers' control**. The process governing the exposures arguably **resembles randomized experiments**.

RCT

- 1 Assignment of treatment is randomized by us
- 2 Treatment is under control by us
- 3 Primary data

Natural Experiment

- 1 Assignment of treatment is randomized by nature
- 2 Treatment is not controlled by us
- 3 Secondary data

Section 2

Regression Discontinuity Design

What is an RDD

- A **regression discontinuity design (RDD)** is a natural-experimental design that aims to determine the causal effects of interventions by identifying a **cutoff** around which an intervention is as if randomized across individuals.

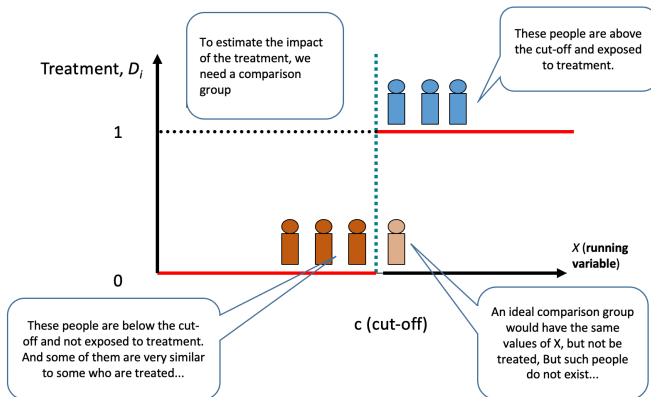


Figure 1: Visual illustration of RDD

Motivating Example

Business objective: What is the causal effect of receiving a Master's degree with Distinction versus Merit on students' future salary?

- Can we run the following simple linear regression and obtain the causal effects?

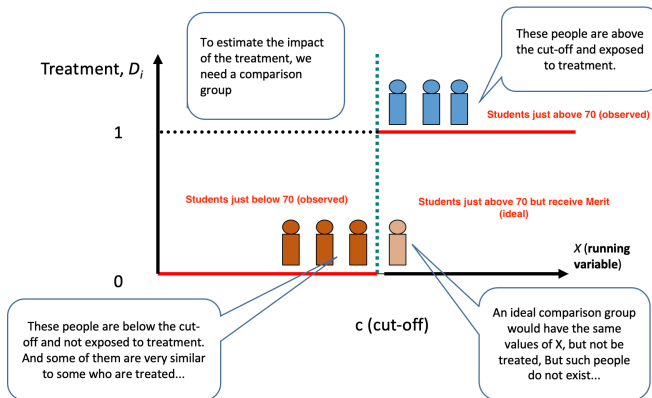
$$Salary_i = \alpha + \beta Distinction_i + \epsilon_i$$

- Can we use RCT?
- Can we use instrumental variables?

A Natural Experiment in the UK

- In the UK Education system, students receiving 70% or above final average grades will receive Distinction while students below 70% will receive Merit.
- The above setting gives us a nice natural experiment:
 - Students may improve their average grades from 60% to 69% by working harder, but they cannot perfectly control their average grades around the cutoff, say from 69.9% to 70%.

Visual Illustration of RDD: An Example of Distinction on Salary



Why RDD Gives Causal Effects?

- For students just above 70%, to measure the treatment effects of receiving Distinction, we would need their *counterfactual salaries if they had not received Distinction*.
- At the same time, because the “running variable” **cannot be perfectly controlled** by the individuals **around the cutoff point**, it’s as if the treatment was randomized near the cutoff. Thus, individuals near the cutoff should be very similar, such that there should be no systematic differences across the treatment and control group.
 - Similar to RCT, we overcome the fundamental problem of causal inference using students just below 70 as the control group.
- All else being equal, a sudden change in the outcome variable at the cutoff can only be attributed to the treatment effect.

When Can We Use an RDD

- An RDD design arises when treatment is assigned based on whether an underlying **continuous variable** crosses a cutoff.
 - The continuous variable is often referred to as the **running variable**.
- **AND** the characteristic cannot be perfectly manipulated by individuals
 - We should only focus on individuals close to the cutoff point.

Exercise: eBay endorses sellers with 10,000 orders as Gold Seller. Can we use RDD to identify the causal effect of receiving Gold Seller endorsement on seller sales?

Section 3

Implementation of RDD

Step 1: Select Sample of Analysis

- ① Determine the bandwidth above and below the cutoff and select the subset of individuals within the bandwidth
 - e.g., if we choose a bandwidth of 0.5, we need to filter out students with average scores between 69.5 and 70.5
- We face a trade-off when selecting the bandwidth: If we choose a smaller bandwidth around the cut-off
 - Pros: Individuals should be more similar around the cutoff, thus it is more likely the control group and treatment group are “as-if randomized”, thus higher internal validity.
 - Cons: We have a smaller subset of subjects which may not be representative of remaining individuals, thus lower external validity; We have a smaller sample size due to fewer individuals selected
- In practice, there is no specific rule how to determine the bandwidth. We need to run a set of different bandwidths as **robustness checks**.

Step 2: Examine Continuity of Observed Characteristics

- ② Examine if other characteristics of the treatment group and control group are continuous at the cut-off point.
 - The idea is similar to “randomization check” in an RCT.

Step 3: Data Analysis

- 3 Regress the outcome variable on the treatment indicator to obtain the causal effect.

$$Y_i = \beta_0 + \beta_1 Treated + \epsilon_i$$

- *Treated* is a binary variable for whether or not the running variable is above the cutoff.
- Sometimes, we may also want to control the running variable in the regression to mitigate its confounding effects.

The Causal Effect of Distinction on Salary

- Generate a synthetic dataset

```
1  pacman::p_load(dplyr,fixest,modelsummary)
2  n <- 1000 # 1000 individuals
3  set.seed(888)
4  score<-runif(n,61,75) # generate scores between 61 and 75
5  experience<-runif(n,0,3) # generate work experience between 0 and 3
6
7  salary<-30000+ 2000*(score>=70) + # causal effect is 2000
8    500 * score + 400*experience + rnorm(n,0,800)
9
10 data_rdd <- data.frame(ID = 1:n,
11                          score = score,
12                          experience = experience,
13                          salary = salary)
```

- Generate the treatment indicator, Distinction, in the dataset using dplyr

Linear Regression Analyses

- Run a linear regression: `salary ~ Distinction`

(1)	
(Intercept)	63 306.835*** (57.722)
Distinction	5533.565*** (94.638)
Num.Obs.	1000
R2	0.774

+ $p < 0.1$, * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$

- The result suggests that Distinction can increase the salary by 5.5k, which is far from the true causal effect.

RDD Analysis

- Step 1: Select a bandwidth around the cutoff, between 68% to 72%
- Step 2: Examine discontinuity of other variables (randomization check).
- Step 3: Run the linear regression on the smaller sample.

(1)	
(Intercept)	65 201.101*** (81.291)
Distinction	2898.285*** (110.725)
Num.Obs.	282
R2	0.710
+ p < 0.1, * p < 0.05, ** p < 0.01, *** p < 0.001	


- Let's try other bandwidths on the Quarto document. As we tighten the bandwidth, what do you find?

Visualization of RDD

- rdrobust package provides a nice visualization tool for RDD.

```
1 pacman::p_load(rdrobust)
2 rdplot(y = salary, # outcome variable
3       x = score, # x is the running variable
4       c = 70, # c is the cutoff point
5       p = 2 # polynomial order to fit the trends
6 )
```

Visualization of RDD



Regression Discontinuity in Time

- A natural experiment occurred on a day affecting all customers, we can then implement a **Regression Discontinuity in Time design (RDiT)** as follows
 - The running variable is time; the cutoff is the day on which the natural experiment took place

$$Y_{it} = \alpha + \beta_1 Post_{it} + \mu_{it}$$

- The underlying assumption for RDiT is that, the pre-treatment outcomes before the natural experiment are good counterfactuals for the post-treatment outcomes if the natural experiment had not happened.
 - For the underlying assumption to hold, we need to take a short time window before and after the natural experiment, e.g., 7 days, 14 days, or 30 days.
- The coefficient β_1 then measures the changes in the outcome variable before and after the natural experiment.

After-class Reading

- (recommended) [Quasi-experiment](#) (Econometrics with R)